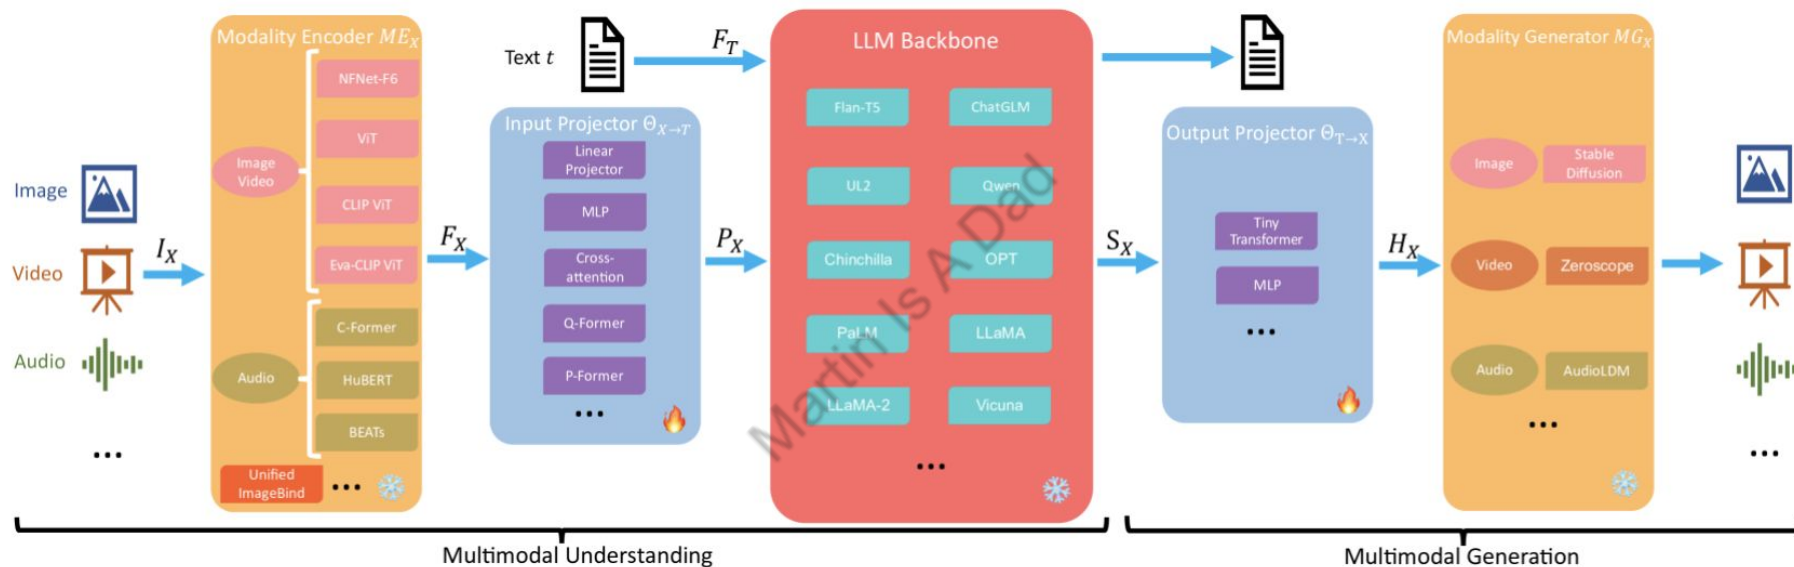


MultiModal LLM Key Components

Component	Modality Encoders	Input Projector	LLM Backbone	Output Projector	Modality Generator
Primary Function	Transform raw data from various modalities (image, audio, 3D, etc.) into numerical feature representations (embeddings).	Align encoded features from different modalities into a common space, typically compatible with the LLM backbone's input. Typically using MLPs, cross attentions, Q-Formers etc.	Serves as the central reasoning and language processing engine, integrating aligned multimodal information. Usually pretrained LLMs, say GPT, Gemini, LLaMA etc.	Maps LLM outputs (e.g., signal tokens for generation) into features suitable for modality-specific generators. Typically MLPs or Transformers.	Produces outputs in non-textual modalities (e.g., images, audio) based on input from the Output Projector. This can be StableDiffusion, AudioLDM etc

MultiModal LLM Key Components



From: [\[arxiv 2401.13601v1\]](https://arxiv.org/abs/2401.13601v1) MM-LLMs: Recent Advances in MultiModal Large Language Models

LLaVA: GPT-assisted Visual Instruction Generation

Context type 1: Captions

A group of people standing outside of a black vehicle with various luggage.

Luggage surrounds a vehicle in an underground parking area

People try to fit all of their luggage in an SUV.

The sport utility vehicle is parked in the public garage, being packed for a trip

Some people with luggage near a van that is transporting it.



Context type 2: Boxes

person: [0.681, 0.242, 0.774, 0.694], backpack: [0.384, 0.696, 0.485, 0.914], suitcase: ...<omitted>

Response type 1: conversation

Question: What type of vehicle is featured in the image?

Answer: The image features a black sport utility vehicle (SUV) ...<omitted>

Response type 2: detailed description

The image is an underground parking area with a black sport utility vehicle (SUV) parked. There are three people in the scene, with one person standing closer to the left side of the vehicle, another person in the middle, and the third person on the right side. They are all working together to pack their luggage into the SUV for a trip. ...<omitted>

Response type 3: complex reasoning

Question: What challenges do these people face?

Answer: In the image, a group of people is standing outside a black SUV in a parking area, surrounded by various pieces of luggage, including suitcases and backpacks. They are facing the challenge of fitting all their luggage into the black SUV. There are multiple suitcases and backpacks to be packed, which suggests that the group has a significant amount of belongings ...<omitted>

LLaVA: Architecture




Figure 1: LLaVA network architecture.

From: [\[arxiv 2304.08485\]](https://arxiv.org/abs/2304.08485) Visual Instruction Tuning

BLIP-2: Overview


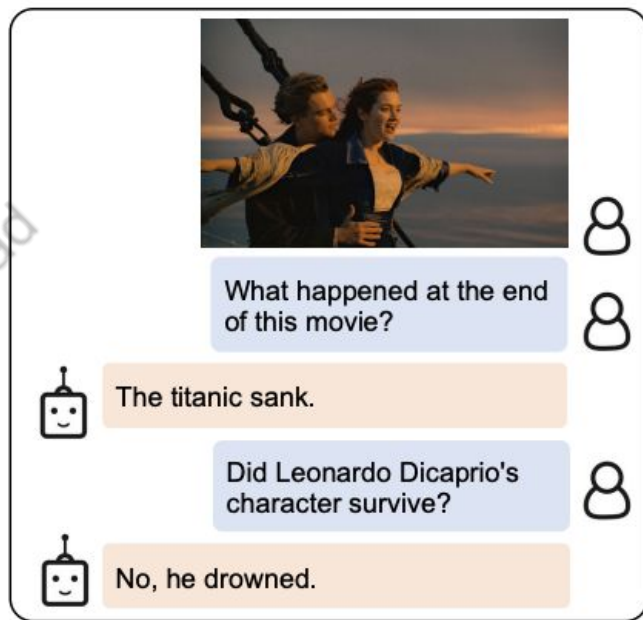

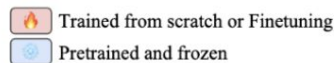



Figure 1. Overview of BLIP-2's framework. We pre-train a lightweight Querying Transformer following a two-stage strategy to bridge the modality gap. The first stage bootstraps vision-language representation learning from a frozen image encoder. The second stage bootstraps vision-to-language generative learning from a frozen LLM, which enables zero-shot instructed image-to-text generation (see Figure 4 for more examples).



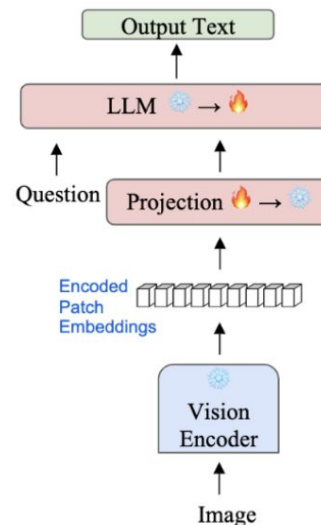
MLLM Training Comparison



Flamingo



BLIP-2 /
InstructBLIP



LLaVA

Current LLM's Fake Reasoning




Figure 26. A simplified attribution graph of Haiku adding two-digit numbers. Features of the inputs feed into separable processing pathways.

[View detailed graph](#)

Human: Answer in one word. What is 36+59?

Assistant: 95

Human: Briefly, how did you get that?

Assistant: I added the ones (6+9=15), carried the 1, then added the tens (3+5+1=9), resulting in 95.

Future

